


# Learning to Sample in an Online Fashion for Robot Motion Planning

Rogério Bonatti, Ratnesh Madaan,  
Brian Okorn, Sam Zeng



# What are Rapidly-Exploring Random Trees (RRT)?



RRT: The Piano-Movers Problem

<https://www.youtube.com/watch?v=rPgZyq15Z-Q&>

# What are Rapidly-Exploring Random Trees?




RRT: The Piano-Movers Problem

<https://www.youtube.com/watch?v=rPgZyq15Z-Q&>


## Related work:

### Learning heuristics



2017 Bhardwaj, Mohak, Sanjiban Choudhury, and Sebastian Scherer use a clairvoyant oracle to **learn heuristics** for search-based planners


### Learning distribution bias




Brian Ichter, James Harrison, and Marco Pavone, "Learning Sampling Distributions for Robot Motion Planning." *arXiv preprint* 2017

# Contributions

Where should I sample here?



What about now?



- Learn a **sampling distribution** using a representation **jointly** over the
  - fully-observable **environment** and
  - the **current state of the search tree**




What is a good sample?

# Supervised sampling

Sample N valid states

Rank by fitness

Pick best K samples



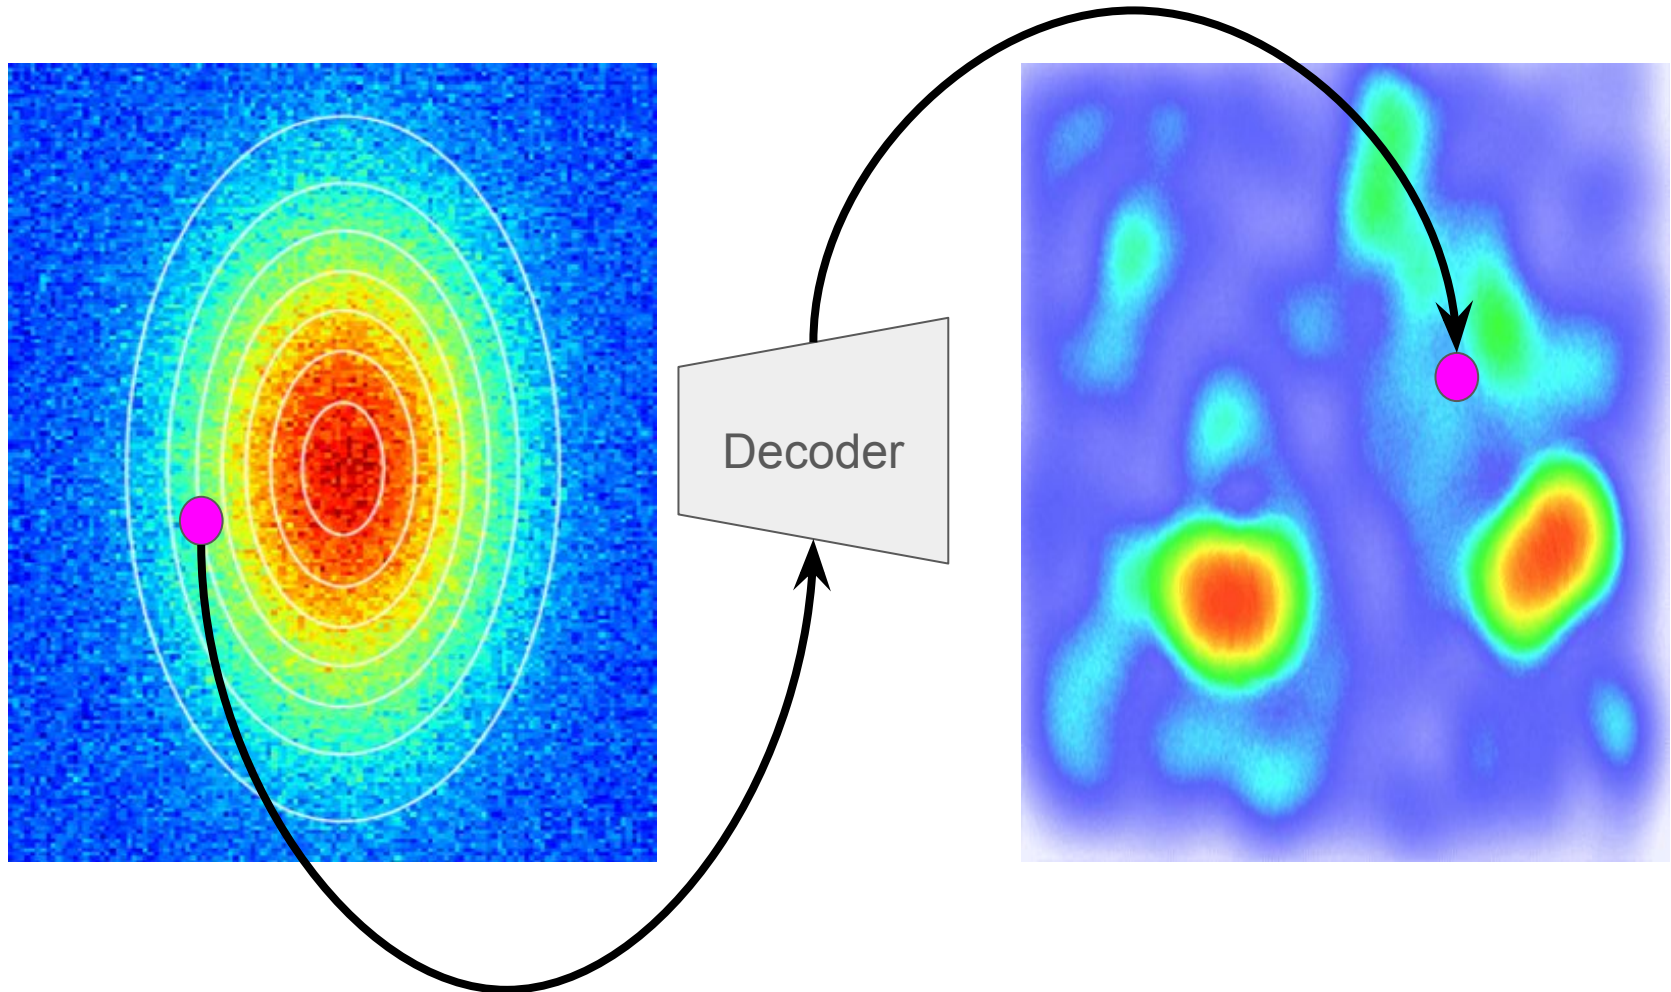
How do we mimic this distribution?




# Variational auto-encoder

Latent space:  $\mathbf{z} \in \mathbb{R}^n$

Sample space:  $\mathbf{x} \in \mathbb{R}^2$




# RRT Sampling CVAE Network

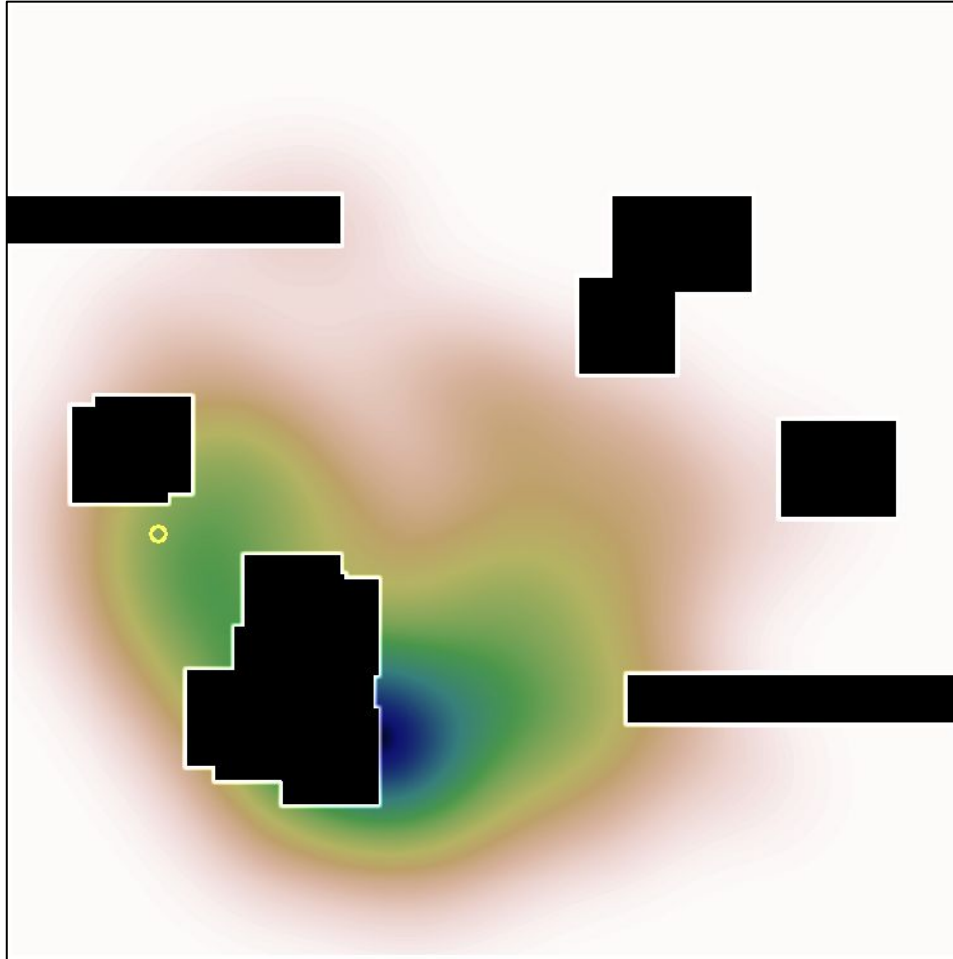


# Results: supervised training


Sequence of test time



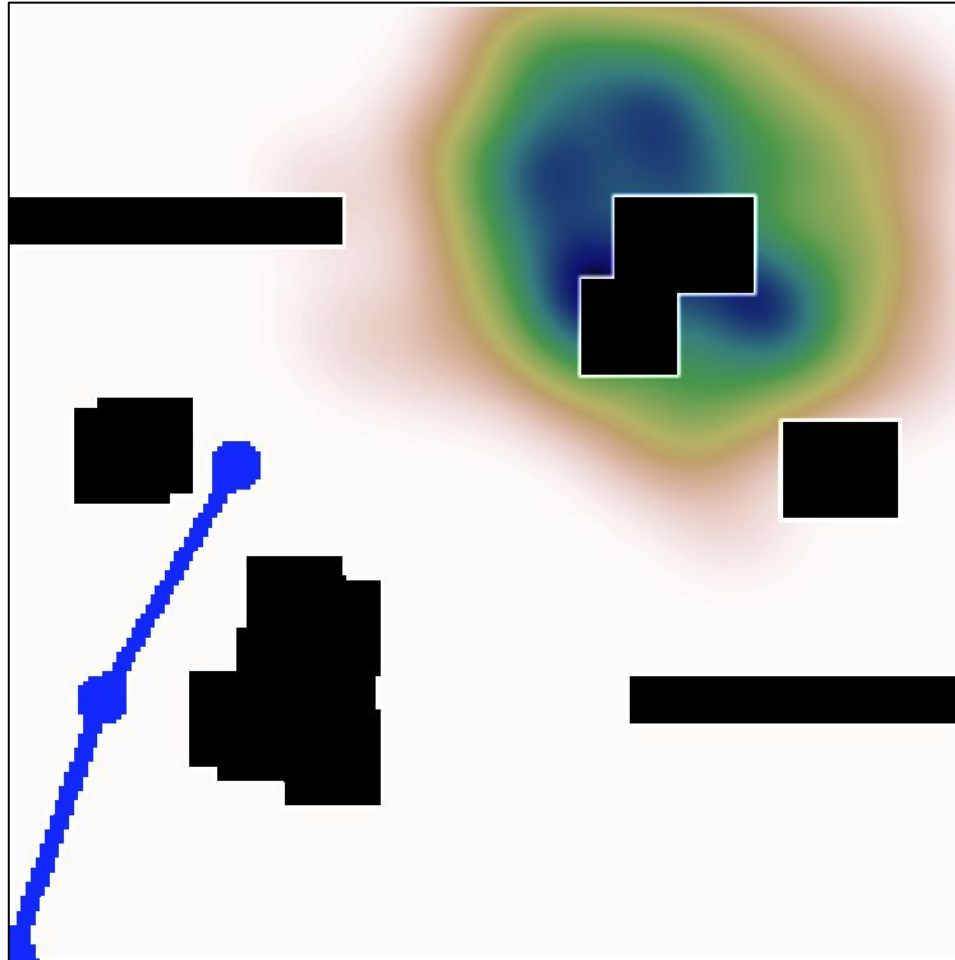
# Solving the planning problem



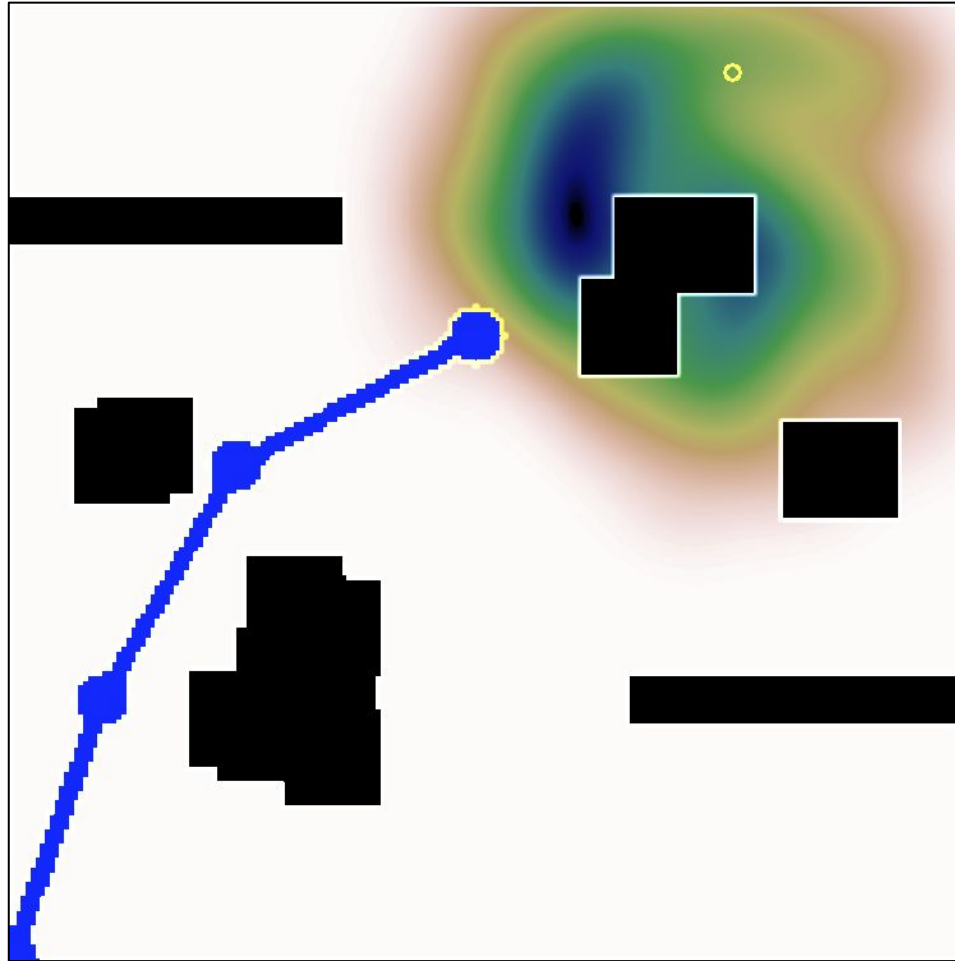
# Solving the planning problem




# Solving the planning problem



# Solving the planning problem






# Solving the planning problem






# Example of distributions in other environments




# DAgger Training

RRT tree




Our trees




# DAgger results

RRT tree




DAgger tree (7 iterations)




What about harder problems?

# Extension to higher state-spaces: 5DOF manipulator


Start



Samples



Goal

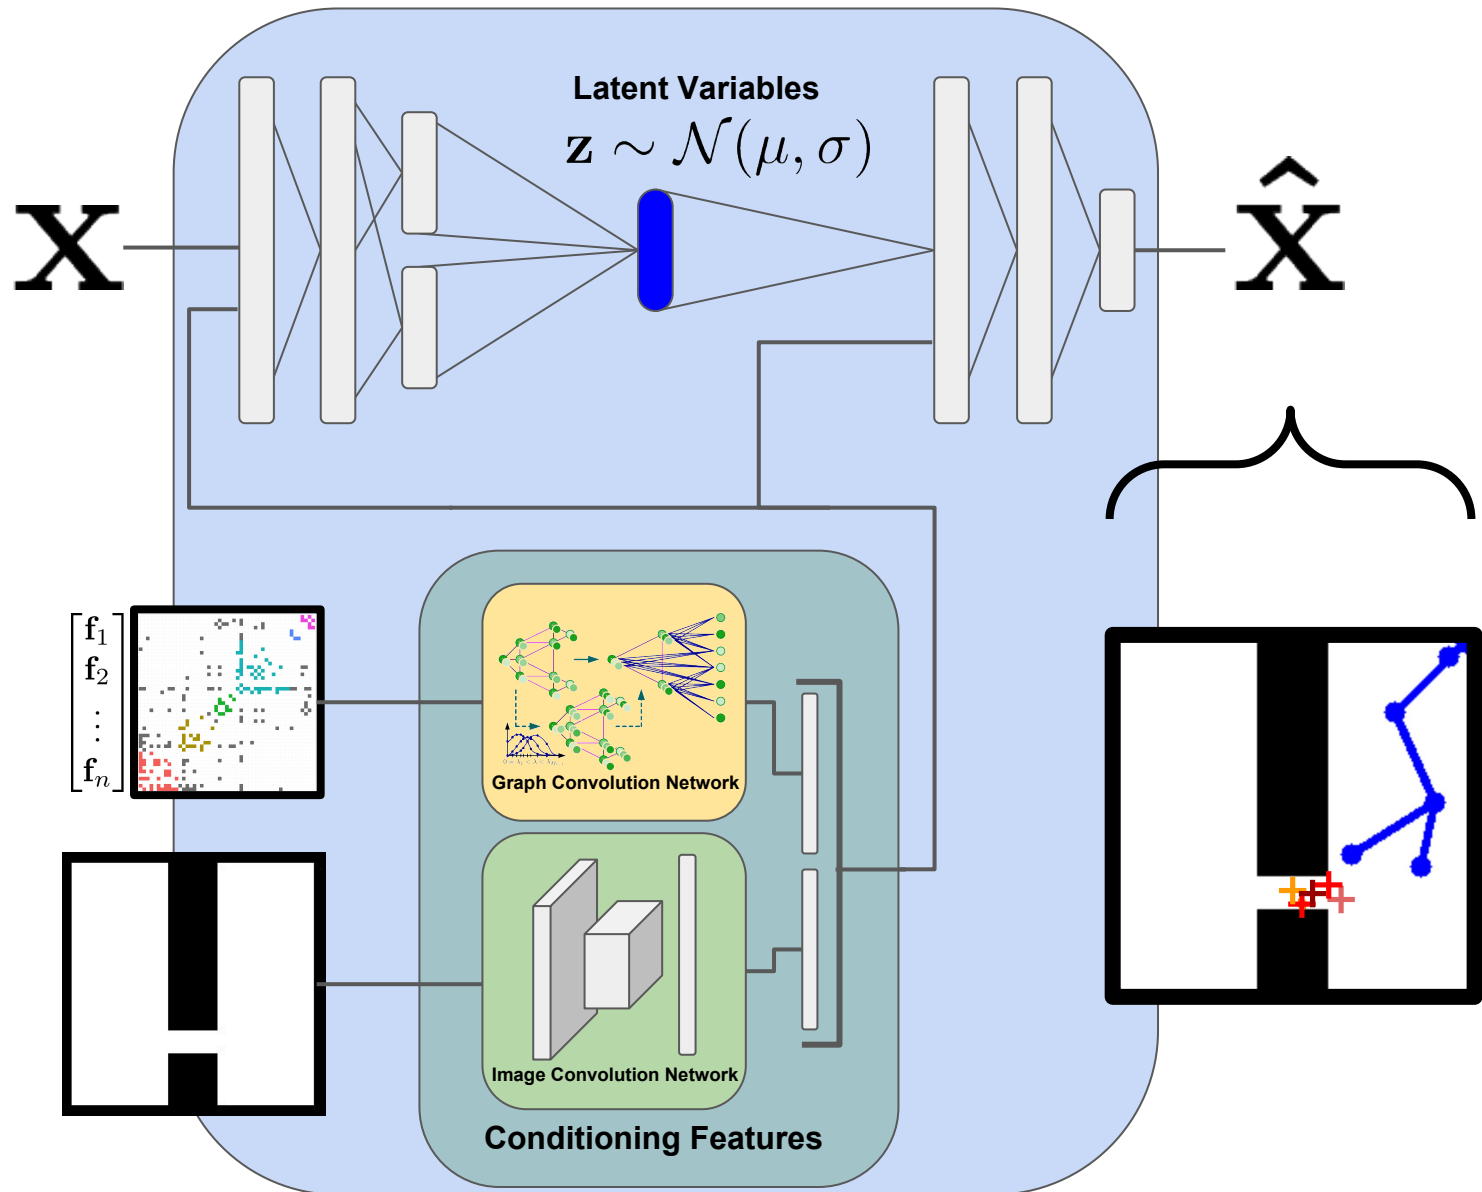


Differences with respect to current approach (2D point robot):

- Current search tree is passed to CVAE as graph instead of an image
  - Adjacency map
  - Featurization (values of each state)
- Cost-to-go function is obtained by solving an optimal planner from sample to goal


Currently being implemented in OMPL for final report

# RRT Sampling With Graph CVAE Network

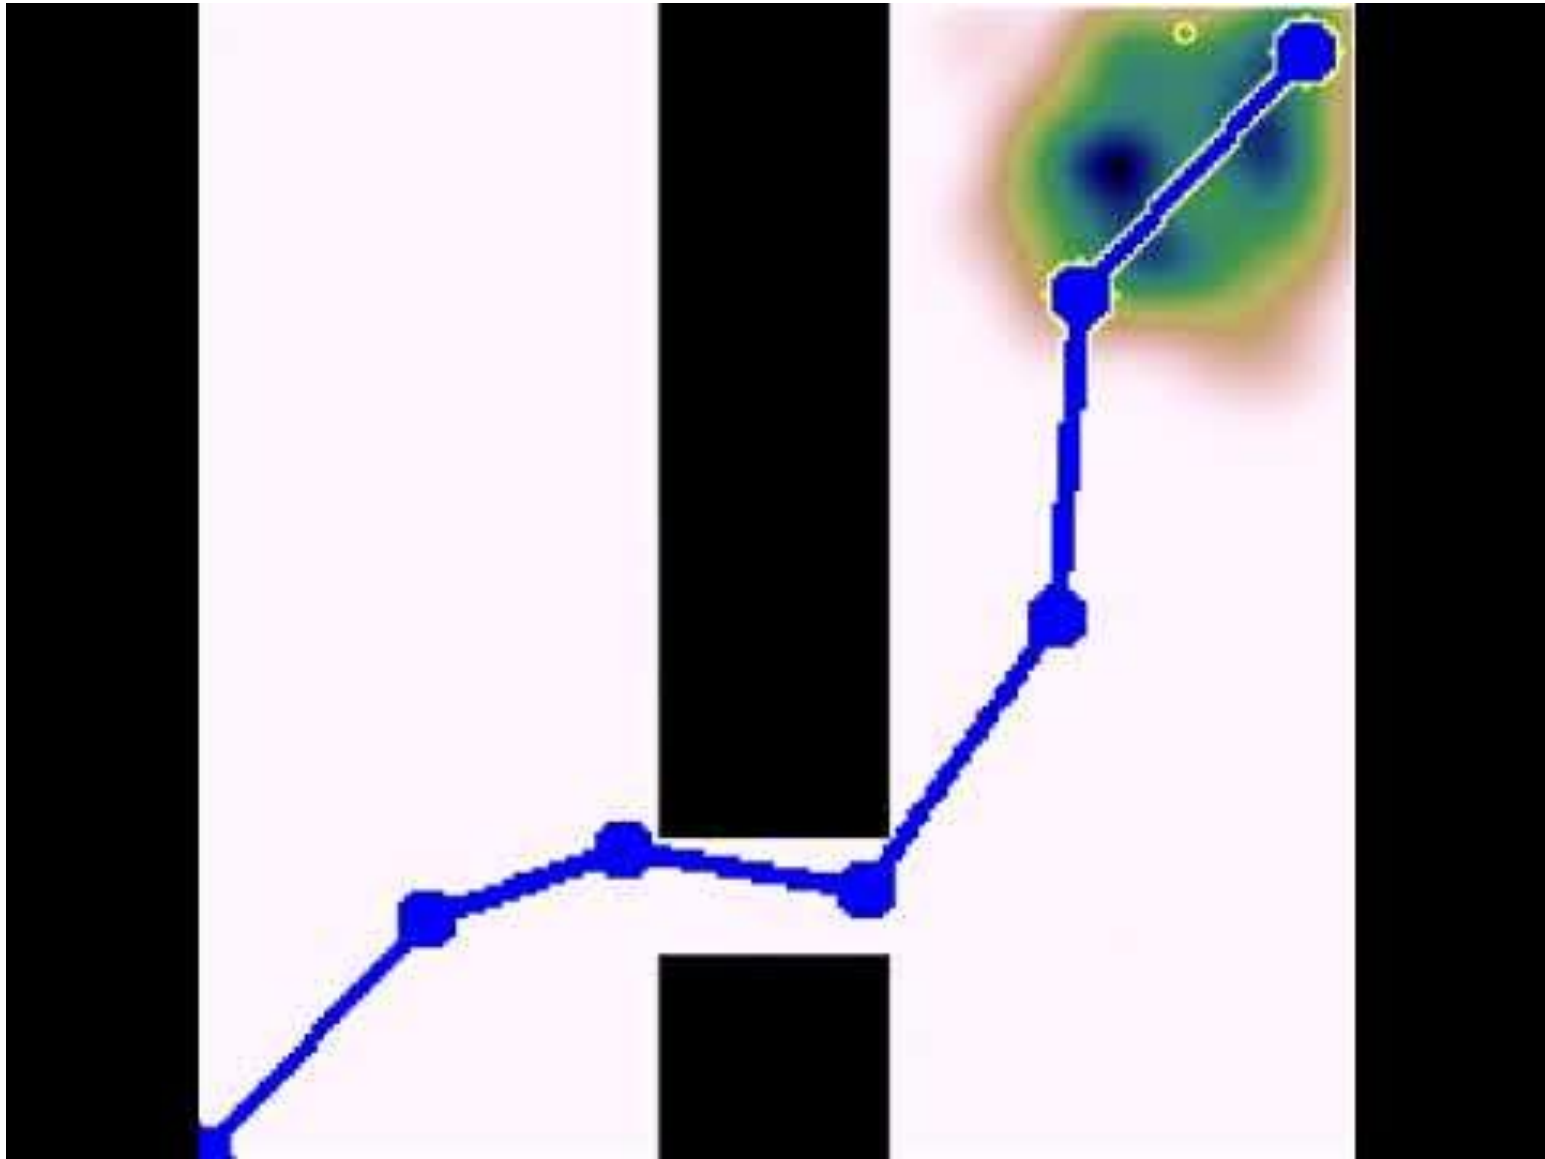


# Questions

- RRT Planning
- CVAE Sampling
- Convolutional Sampler
- DAgger Training
- Higher Dimensions
- Graph Sampler

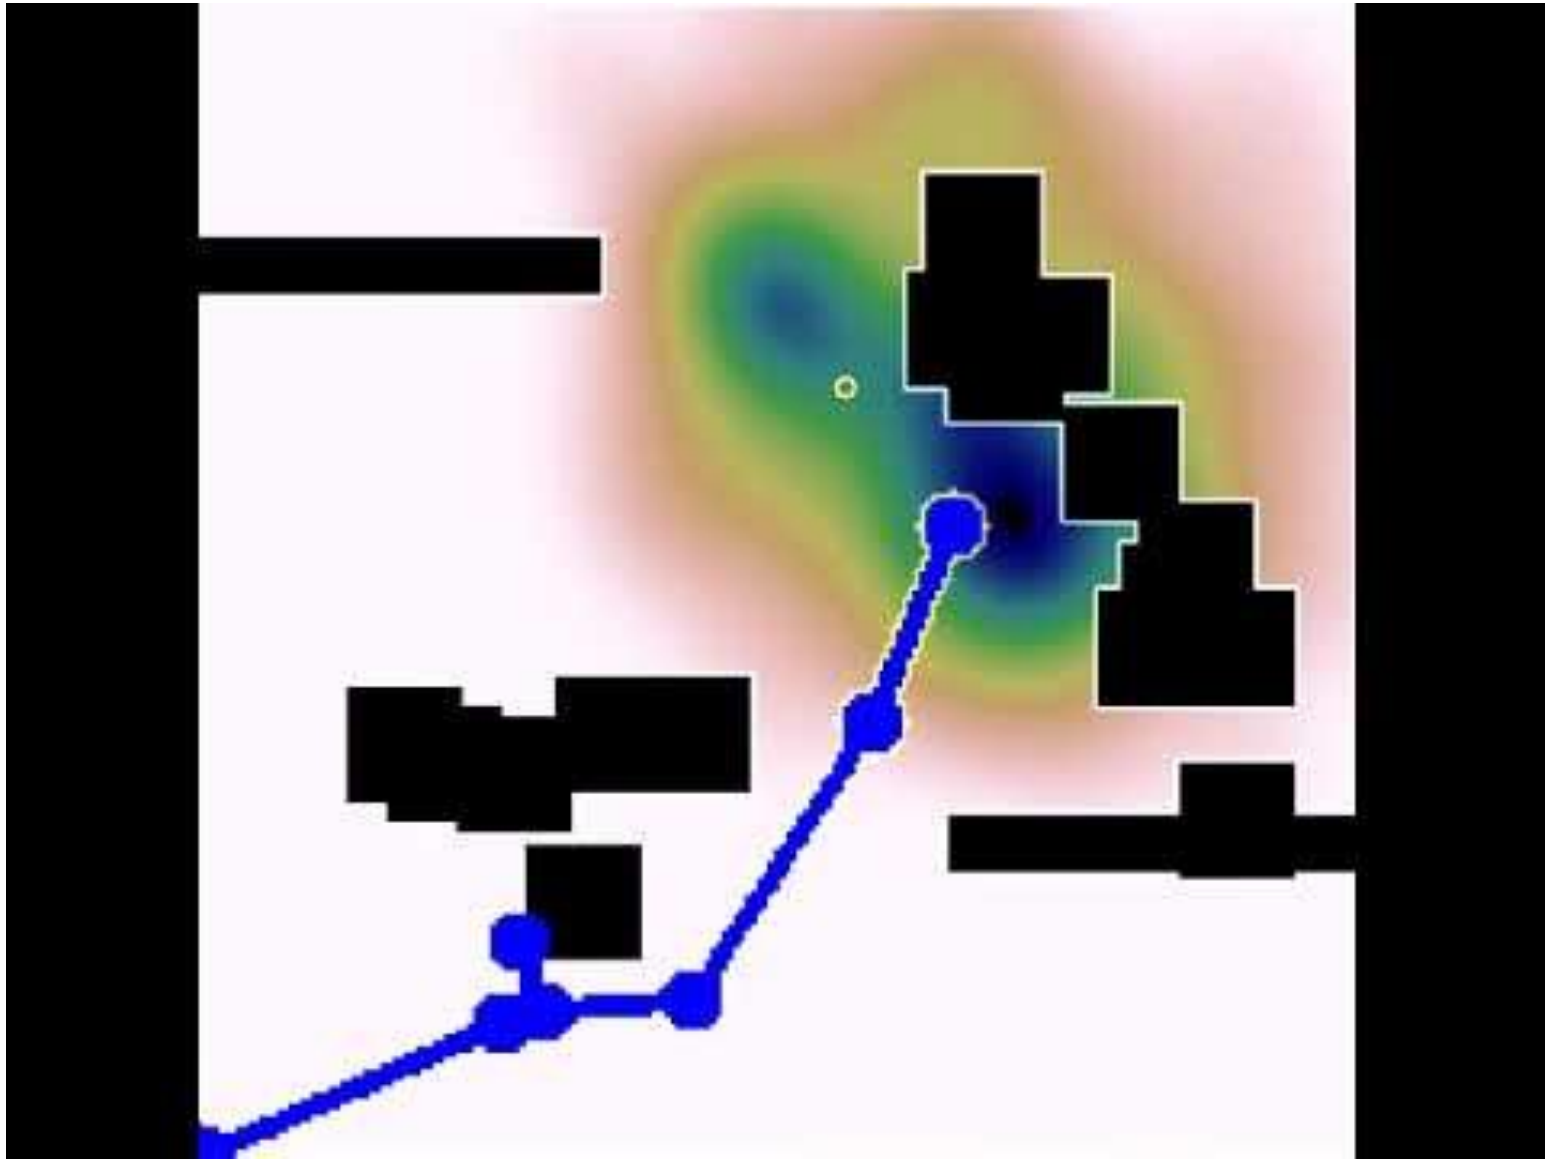


# Shifting Gaps







# Forest




# Mazes




# Single Bugtrap



# Multiple Bugtraps



# Gaps and Forest





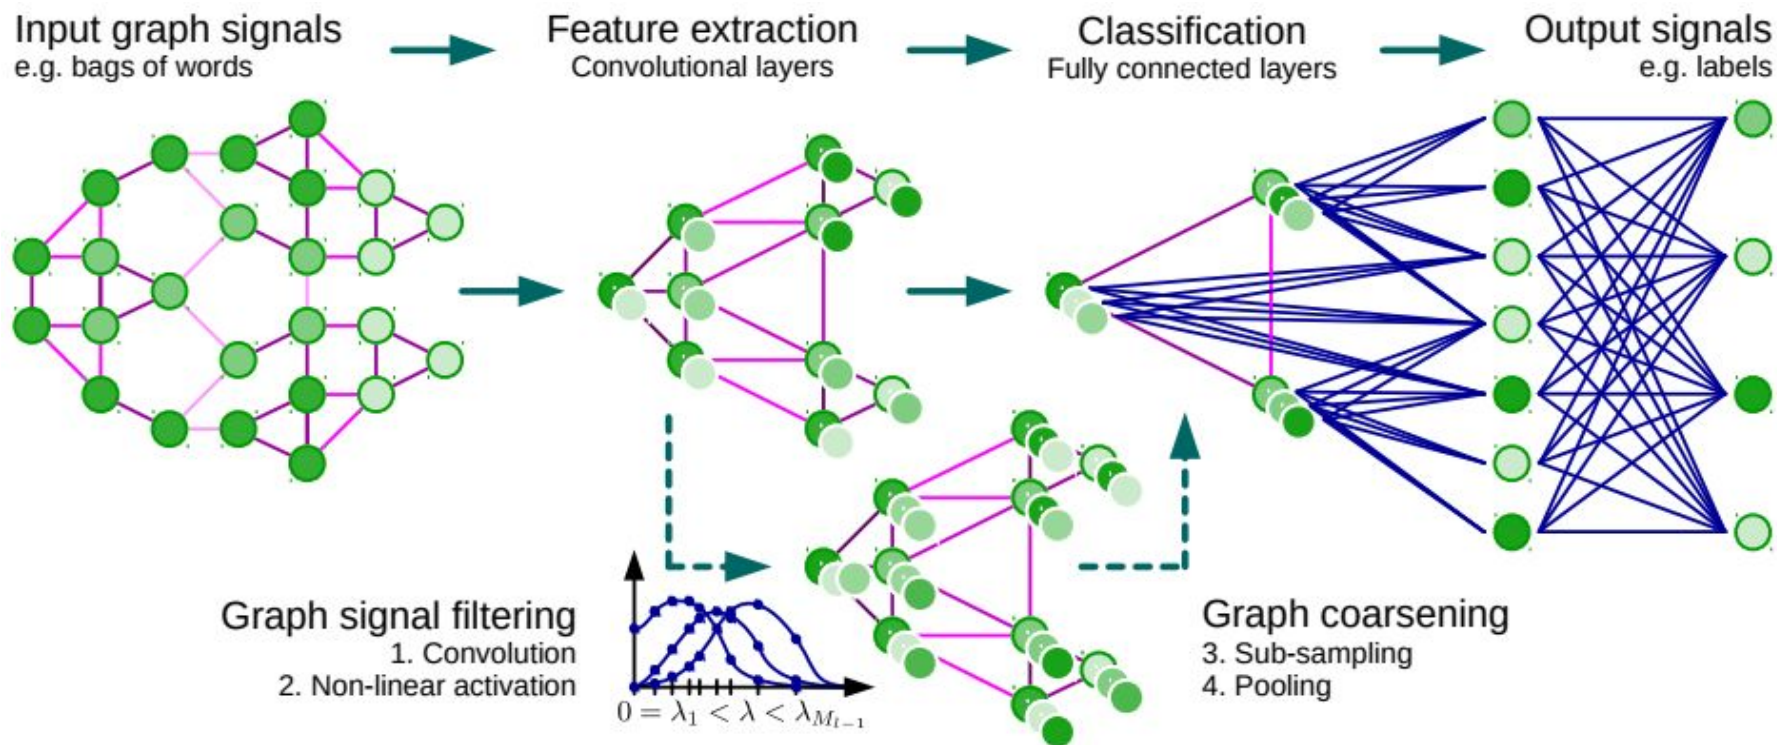
**BACKUP**

## Some samples are more optimistic-in-the-face-of-partial-observability than others


- Ideally, we would want the learnt policy to take risks aka be “greedy” sometimes, depending on the planning progress, (number of (failed)) collision checks, etc.
- For instance, a simple conditioning variable could be the number of collision checks failed till now. If this is high, ideally the distribution should have more density near the graph.  
On the contrary, if a lot of collisions were happening, the distribution can take risks
- A simple #gupta way of encoding this is add another channel with circles on the states that failed collision checks  
(it's kinda weird - one channel has black for obstacles - a bad region indicator, one channel has black for the graph - sample close to this indicator, one channel has black for the failed nodes - sample away from this indicator.  
I guess we should have some color for good things, and bad things?




# Graph-based Features: Localized Spectral Filtering



# Graph Coarsening using Heavy Edge Matching



# What are Rapidly-Exploring Random Trees?



RRT: The Piano-Movers Problem

<https://www.youtube.com/watch?v=rPgZyq15Z-Q&>